

Satya Sudheer Yalamarthy

satya7sudheer@gmail.com | +91 9535760123 | Bengaluru, India | linkedin.com/in/satyasudheer

August 27, 2026

Hiring Team

GitLab

India (Remote)

Re: Director of Engineering, Organizations & Cells

Dear Hiring Team,

I'm writing to apply for the Director of Engineering role leading Organizations and Cells. The charter you've described — cell provisioning and lifecycle, cohort-based tenant migration with zero data loss and verified rollback, a multi-cloud and multi-region footprint — is the same class of problem I've spent the last several years building and leading through, at a smaller but structurally similar scale, and I'd welcome the chance to bring that experience to a program of this size.

At TransUnion, I architected a multi-tenant, multi-region control plane for real-time IP, device, and identity reputation, sustaining ~1,000 QPS at 99.9%+ availability while designing explicitly for blast-radius containment across regions — migrating the platform from batch enrichment to event-driven streaming on Kafka and Flink. At Paytm, I led a settlement-processing migration that re-partitioned microservice and database boundaries in a phased, verified-rollback rollout, delivering a 4x increase in daily transaction throughput and a 29-30% infrastructure cost reduction as the footprint grew — the same 'sequencing across many teams, not writing the code' problem this role is built around.

I've led engineering organizations directly (~20-26 engineers) and cross-functionally through managers (~25 engineers across cities), setting the technical bar and operating cadence across data engineering, SRE, and product, and hiring and developing the managers who ran those teams day to day. Cloud infrastructure depth spans AWS, GCP, and Azure, including AWS Solutions Architect, Developer, and SysOps certifications, Azure Fundamentals, and Google's Professional Cloud Architect certification.

On the AI-native side, I designed and built an agentic, Retrieval-Augmented Generation system at TransUnion — a six-stage pipeline using a pgvector store and the Anthropic SDK behind a human-in-the-loop approval workflow — and had to reason directly about what agentic access does to read volume, concurrency, and authorization before it went into production, rather than retrofitting it afterward. That's a design discipline I'd bring to how Cells expose data to AI-native workloads.

GitLab's remote, asynchronous, documentation-first way of working is one I've operated in for years across distributed teams and time zones, and the scale and technical weight of this charter — rebuilding GitLab.com's substrate without customers ever seeing the seam — is exactly the kind of high-consequence infrastructure work I want to be doing next. I'd welcome the opportunity to discuss how my background fits this role.

Thank you for your time and consideration.

Sincerely,
Satya Sudheer Yalamarthy